

```
In [ ]: Name: N.S.Sai Parasanth  
        Regno: 21BCE8305  
        Date: 22nd septmber, 2023
```

```
In [1]: import numpy as np  
        import pandas as pd  
        import matplotlib.pyplot as plt  
        import seaborn as sns
```

```
In [2]: data=pd.read_csv("Employee-Attrition.csv")
```

```
In [3]: data.head()
```

```
Out[3]:
```

	Age	Attrition	BusinessTravel	DailyRate	Department	DistanceFromHome	Education	Educ
0	41	Yes	Travel_Rarely	1102	Sales	1	2	Lif
1	49	No	Travel_Frequently	279	Research & Development	8	1	Lif
2	37	Yes	Travel_Rarely	1373	Research & Development	2	2	
3	33	No	Travel_Frequently	1392	Research & Development	3	4	Lif
4	27	No	Travel_Rarely	591	Research & Development	2	1	

5 rows × 35 columns

```
In [4]: data.tail()
```

```
Out[4]:
```

	Age	Attrition	BusinessTravel	DailyRate	Department	DistanceFromHome	Education	E
1465	36	No	Travel_Frequently	884	Research & Development	23	2	
1466	39	No	Travel_Rarely	613	Research & Development	6	1	
1467	27	No	Travel_Rarely	155	Research & Development	4	3	
1468	49	No	Travel_Frequently	1023	Sales	2	3	
1469	34	No	Travel_Rarely	628	Research & Development	8	3	

5 rows × 35 columns

```
In [5]: data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1470 entries, 0 to 1469
Data columns (total 35 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Age                                   1470 non-null   int64
1   Attrition                           1470 non-null   object
2   BusinessTravel                      1470 non-null   object
3   DailyRate                           1470 non-null   int64
4   Department                          1470 non-null   object
5   DistanceFromHome                    1470 non-null   int64
6   Education                           1470 non-null   int64
7   EducationField                      1470 non-null   object
8   EmployeeCount                       1470 non-null   int64
9   EmployeeNumber                      1470 non-null   int64
10  EnvironmentSatisfaction              1470 non-null   int64
11  Gender                              1470 non-null   object
12  HourlyRate                          1470 non-null   int64
13  JobInvolvement                      1470 non-null   int64
14  JobLevel                            1470 non-null   int64
15  JobRole                             1470 non-null   object
16  JobSatisfaction                     1470 non-null   int64
17  MaritalStatus                      1470 non-null   object
18  MonthlyIncome                      1470 non-null   int64
19  MonthlyRate                         1470 non-null   int64
20  NumCompaniesWorked                  1470 non-null   int64
21  Over18                             1470 non-null   object
22  OverTime                            1470 non-null   object
23  PercentSalaryHike                   1470 non-null   int64
24  PerformanceRating                   1470 non-null   int64
25  RelationshipSatisfaction             1470 non-null   int64
26  StandardHours                       1470 non-null   int64
27  StockOptionLevel                   1470 non-null   int64
28  TotalWorkingYears                   1470 non-null   int64
29  TrainingTimesLastYear               1470 non-null   int64
30  WorkLifeBalance                     1470 non-null   int64
31  YearsAtCompany                      1470 non-null   int64
32  YearsInCurrentRole                  1470 non-null   int64
33  YearsSinceLastPromotion              1470 non-null   int64
34  YearsWithCurrManager                1470 non-null   int64
dtypes: int64(26), object(9)
memory usage: 402.1+ KB
```

```
In [6]: data.describe()
```

```
Out[6]:
```

	Age	DailyRate	DistanceFromHome	Education	EmployeeCount	EmployeeNui
count	1470.000000	1470.000000	1470.000000	1470.000000	1470.0	1470.00
mean	36.923810	802.485714	9.192517	2.912925	1.0	1024.86
std	9.135373	403.509100	8.106864	1.024165	0.0	602.02
min	18.000000	102.000000	1.000000	1.000000	1.0	1.00
25%	30.000000	465.000000	2.000000	2.000000	1.0	491.25
50%	36.000000	802.000000	7.000000	3.000000	1.0	1020.50
75%	43.000000	1157.000000	14.000000	4.000000	1.0	1555.75
max	60.000000	1499.000000	29.000000	5.000000	1.0	2068.00

8 rows × 26 columns

Handling the null values

```
In [7]: data.isnull().any()
```

```
Out[7]: Age                False
Attrition                 False
BusinessTravel            False
DailyRate                 False
Department                False
DistanceFromHome          False
Education                 False
EducationField             False
EmployeeCount              False
EmployeeNumber             False
EnvironmentSatisfaction   False
Gender                    False
HourlyRate                 False
JobInvolvement             False
JobLevel                  False
JobRole                   False
JobSatisfaction            False
MaritalStatus              False
MonthlyIncome              False
MonthlyRate                False
NumCompaniesWorked         False
Over18                    False
OverTime                   False
PercentSalaryHike          False
PerformanceRating          False
RelationshipSatisfaction   False
StandardHours              False
StockOptionLevel           False
TotalWorkingYears          False
TrainingTimesLastYear      False
WorkLifeBalance            False
YearsAtCompany             False
YearsInCurrentRole         False
YearsSinceLastPromotion    False
YearsWithCurrManager        False
dtype: bool
```

```
In [8]: data.isnull().sum()
```


```
Out[8]: Age                                0
Attrition                                0
BusinessTravel                           0
DailyRate                                0
Department                               0
DistanceFromHome                         0
Education                                0
EducationField                            0
EmployeeCount                             0
EmployeeNumber                           0
EnvironmentSatisfaction                   0
Gender                                    0
HourlyRate                                0
JobInvolvement                           0
JobLevel                                 0
JobRole                                  0
JobSatisfaction                           0
MaritalStatus                            0
MonthlyIncome                             0
MonthlyRate                              0
NumCompaniesWorked                       0
Over18                                    0
OverTime                                  0
PercentSalaryHike                         0
PerformanceRating                         0
RelationshipSatisfaction                   0
StandardHours                             0
StockOptionLevel                          0
TotalWorkingYears                        0
TrainingTimesLastYear                     0
WorkLifeBalance                           0
YearsAtCompany                             0
YearsInCurrentRole                         0
YearsSinceLastPromotion                   0
YearsWithCurrManager                      0
dtype: int64
```

```
In [9]: cor=data.corr()
```

```
C:\Users\Prasanth Nimmala\AppData\Local\Temp\ipykernel_8884\1426905697.py:1:
FutureWarning: The default value of numeric_only in DataFrame.corr is depreca
ted. In a future version, it will default to False. Select only valid columns
or specify the value of numeric_only to silence this warning.
  cor=data.corr()
```

```
In [10]: fig=plt.figure(figsize=(18,18))
sns.heatmap(cor,annot=True)
```


Out[10]: <Axes: >



Outliers


```
In [11]: sns.boxplot(data["Age"])
```

```
Out[11]: <Axes: >
```



```
In [12]: sns.boxplot(data["DailyRate"])
```

```
Out[12]: <Axes: >
```



```
In [29]: data.describe()
```

```
Out[29]:
```

	Age	DailyRate	DistanceFromHome	Education	EmployeeCount	EmployeeNui
count	1470.000000	1470.000000	1470.000000	1470.000000	1470.0	1470.00
mean	36.923810	802.485714	9.192517	2.912925	1.0	1024.86
std	9.135373	403.509100	8.106864	1.024165	0.0	602.02
min	18.000000	102.000000	1.000000	1.000000	1.0	1.00
25%	30.000000	465.000000	2.000000	2.000000	1.0	491.25
50%	36.000000	802.000000	7.000000	3.000000	1.0	1020.50
75%	43.000000	1157.000000	14.000000	4.000000	1.0	1555.75
max	60.000000	1499.000000	29.000000	5.000000	1.0	2068.00

8 rows × 26 columns

In [30]: `data.head()`


Out[30]:

	Age	BusinessTravel	DailyRate	Department	DistanceFromHome	Education	EmployeeCount
0	41	Travel_Rarely	1102	Sales	1	2	1
1	49	Travel_Frequently	279	Research & Development	8	1	1
2	37	Travel_Rarely	1373	Research & Development	2	2	1
3	33	Travel_Frequently	1392	Research & Development	3	4	1
4	27	Travel_Rarely	591	Research & Development	2	1	1

5 rows × 33 columns


In [14]: `fig, axes = plt.subplots(2,2)`
`sns.boxplot(data=data["YearsInCurrentRole"],ax=axes[0,0])`
`sns.boxplot(data=data["YearsSinceLastPromotion"],ax=axes[0,1])`
`sns.boxplot(data=data["YearsWithCurrManager"],ax=axes[1,0])`
`sns.boxplot(data=data["WorkLifeBalance"],ax=axes[1,1])`

Out[14]: <Axes: >



```
In [15]: fig, axes = plt.subplots(2,2)
sns.boxplot(data=data["DistanceFromHome"],ax=axes[0,0])
sns.boxplot(data=data["TotalWorkingYears"],ax=axes[0,1])
sns.boxplot(data=data["HourlyRate"],ax=axes[1,0])
sns.boxplot(data=data["YearsAtCompany"],ax=axes[1,1])
```

Out[15]: <Axes: >



Handling the outliers

```
In [16]: YearsInCurrentRole_q1 = data.YearsInCurrentRole.quantile(0.25)
YearsInCurrentRole_q3 = data.YearsInCurrentRole.quantile(0.75)
IQR_YearsInCurrentRole=YearsInCurrentRole_q3-YearsInCurrentRole_q1
upperlimit_YearsInCurrentRole=YearsInCurrentRole_q3+1.5*IQR_YearsInCurrentRole
lower_limit_YearsInCurrentRole =YearsInCurrentRole_q1-1.5*IQR_YearsInCurrentRo
median_YearsInCurrentRole=data["YearsInCurrentRole"].median()
data['YearsInCurrentRole'] = np.where(
    (data['YearsInCurrentRole'] > upperlimit_YearsInCurrentRole),
    median_YearsInCurrentRole,
    data['YearsInCurrentRole']
)
```

```
In [17]: YearsSinceLastPromotion_q1 = data.YearsSinceLastPromotion.quantile(0.25)
YearsSinceLastPromotion_q3 = data.YearsSinceLastPromotion.quantile(0.75)
IQR_YearsSinceLastPromotion=YearsSinceLastPromotion_q3-YearsSinceLastPromotion_q1
upperlimit_YearsSinceLastPromotion=YearsSinceLastPromotion_q3+1.5*IQR_YearsSin
lower_limit_YearsSinceLastPromotion =YearsSinceLastPromotion_q1-1.5*IQR_YearsS
median_YearsSinceLastPromotion=data["YearsSinceLastPromotion"].median()
data['YearsSinceLastPromotion'] = np.where(
    (data['YearsSinceLastPromotion'] > upperlimit_YearsSinceLastPromotion),
    median_YearsSinceLastPromotion,
    data['YearsSinceLastPromotion']
)
```


```
In [18]: YearsWithCurrManager_q1 = data.YearsWithCurrManager.quantile(0.25)
YearsWithCurrManager_q3 = data.YearsWithCurrManager.quantile(0.75)
IQR_YearsWithCurrManager=YearsWithCurrManager_q3-YearsWithCurrManager_q1
upperlimit_YearsWithCurrManager=YearsWithCurrManager_q3+1.5*IQR_YearsWithCurrM
lower_limit_YearsWithCurrManager =YearsWithCurrManager_q1-1.5*IQR_YearsWithCur
median_YearsWithCurrManager=data["YearsWithCurrManager"].median()
data['YearsWithCurrManager'] = np.where(
    (data['YearsWithCurrManager'] > upperlimit_YearsWithCurrManager),
    median_YearsWithCurrManager,
    data['YearsWithCurrManager']
)
```

```
In [19]: TotalWorkingYears_q1 = data.TotalWorkingYears.quantile(0.25)
TotalWorkingYears_q3 = data.TotalWorkingYears.quantile(0.75)
IQR_TotalWorkingYears=TotalWorkingYears_q3-TotalWorkingYears_q1
upperlimit_TotalWorkingYears=TotalWorkingYears_q3+1.5*IQR_TotalWorkingYears
lower_limit_TotalWorkingYears=TotalWorkingYears_q1-1.5*IQR_TotalWorkingYears
median_TotalWorkingYears=data["TotalWorkingYears"].median()
data['TotalWorkingYears'] = np.where(
    (data['TotalWorkingYears'] > upperlimit_TotalWorkingYears),
    median_TotalWorkingYears,
    data['TotalWorkingYears']
)
```

```
In [20]: YearsAtCompany_q1 = data.YearsAtCompany.quantile(0.25)
YearsAtCompany_q3 = data.YearsAtCompany.quantile(0.75)
IQR_YearsAtCompany=YearsAtCompany_q3-YearsAtCompany_q1
upperlimit_YearsAtCompany=YearsAtCompany_q3+1.5*IQR_YearsAtCompany
lower_limit_YearsAtCompany=YearsAtCompany_q1-1.5*IQR_YearsAtCompany
median_YearsAtCompany=data["YearsAtCompany"].median()
data['YearsAtCompany'] = np.where(
    (data['YearsAtCompany'] > upperlimit_YearsAtCompany),
    median_YearsAtCompany,
    data['YearsAtCompany']
)
```

```
In [21]: fig, axes = plt.subplots(2,2)
sns.boxplot(data=data["YearsWithCurrManager"],ax=axes[0,0])
sns.boxplot(data=data["TotalWorkingYears"],ax=axes[0,1])
sns.boxplot(data=data["YearsSinceLastPromotion"],ax=axes[1,0])
sns.boxplot(data=data["YearsAtCompany"],ax=axes[1,1])
```

Out[21]: <Axes: >



```
In [31]: data.head()
```

```
Out[31]:
```

	Age	BusinessTravel	DailyRate	Department	DistanceFromHome	Education	EmployeeCount
0	41	Travel_Rarely	1102	Sales	1	2	1
1	49	Travel_Frequently	279	Research & Development	8	1	1
2	37	Travel_Rarely	1373	Research & Development	2	2	1
3	33	Travel_Frequently	1392	Research & Development	3	4	1
4	27	Travel_Rarely	591	Research & Development	2	1	1

5 rows × 33 columns

```
In [22]: data.drop("EducationField",axis=1,inplace=True)
```

```
In [23]: data.head(2)
```

```
Out[23]:
```

	Age	Attrition	BusinessTravel	DailyRate	Department	DistanceFromHome	Education	Empl
0	41	Yes	Travel_Rarely	1102	Sales	1	2	
1	49	No	Travel_Frequently	279	Research & Development	8	1	

2 rows × 34 columns

```
In [24]: data["BusinessTravel"].unique()
```

```
Out[24]: array(['Travel_Rarely', 'Travel_Frequently', 'Non-Travel'], dtype=object)
```

Splitting the data

```
In [25]: y=data["Attrition"]
```

```
In [26]: y.head()
```

```
Out[26]: 0    Yes
1    No
2    Yes
3    No
4    No
Name: Attrition, dtype: object
```

```
In [27]: data.drop("Attrition",axis=1,inplace=True)
```

```
In [28]: data.head()
```

```
Out[28]:
```

	Age	BusinessTravel	DailyRate	Department	DistanceFromHome	Education	EmployeeCount
0	41	Travel_Rarely	1102	Sales	1	2	1
1	49	Travel_Frequently	279	Research & Development	8	1	1
2	37	Travel_Rarely	1373	Research & Development	2	2	1
3	33	Travel_Frequently	1392	Research & Development	3	4	1
4	27	Travel_Rarely	591	Research & Development	2	1	1

5 rows × 33 columns

Encoding

```
In [32]: from sklearn.preprocessing import LabelEncoder
```

```
In [33]: le=LabelEncoder()
```

```
In [34]: data["BusinessTravel"]=le.fit_transform(data["BusinessTravel"])
```

```
In [35]: data["Department"]=le.fit_transform(data["Department"])
```

```
In [36]: data["Gender"]=le.fit_transform(data["Gender"])
```

```
In [37]: y=le.fit_transform(y)
```

```
In [38]: y
```

```
Out[38]: array([1, 0, 1, ..., 0, 0, 0])
```

```
In [39]: data["JobRole"]=le.fit_transform(data["JobRole"])
```

```
In [40]: data["Over18"]=le.fit_transform(data["Over18"])
```

```
In [41]: data["MaritalStatus"]=le.fit_transform(data["MaritalStatus"])
```

```
In [42]: data["OverTime"]=le.fit_transform(data["OverTime"])
```

In [43]: data.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1470 entries, 0 to 1469
Data columns (total 33 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Age                                    1470 non-null   int64
1   BusinessTravel                        1470 non-null   int32
2   DailyRate                             1470 non-null   int64
3   Department                            1470 non-null   int32
4   DistanceFromHome                      1470 non-null   int64
5   Education                             1470 non-null   int64
6   EmployeeCount                         1470 non-null   int64
7   EmployeeNumber                        1470 non-null   int64
8   EnvironmentSatisfaction               1470 non-null   int64
9   Gender                                1470 non-null   int32
10  HourlyRate                            1470 non-null   int64
11  JobInvolvement                        1470 non-null   int64
12  JobLevel                              1470 non-null   int64
13  JobRole                               1470 non-null   int32
14  JobSatisfaction                       1470 non-null   int64
15  MaritalStatus                         1470 non-null   int32
16  MonthlyIncome                         1470 non-null   int64
17  MonthlyRate                           1470 non-null   int64
18  NumCompaniesWorked                   1470 non-null   int64
19  Over18                                1470 non-null   int32
20  OverTime                              1470 non-null   int32
21  PercentSalaryHike                    1470 non-null   int64
22  PerformanceRating                    1470 non-null   int64
23  RelationshipSatisfaction              1470 non-null   int64
24  StandardHours                        1470 non-null   int64
25  StockOptionLevel                     1470 non-null   int64
26  TotalWorkingYears                    1470 non-null   float64
27  TrainingTimesLastYear                1470 non-null   int64
28  WorkLifeBalance                      1470 non-null   int64
29  YearsAtCompany                       1470 non-null   float64
30  YearsInCurrentRole                   1470 non-null   float64
31  YearsSinceLastPromotion               1470 non-null   float64
32  YearsWithCurrManager                 1470 non-null   float64
dtypes: float64(5), int32(7), int64(21)
memory usage: 338.9 KB
```

train test split

In [44]: `from sklearn.model_selection import train_test_split`
`x_train,x_test,y_train,y_test=train_test_split(data,y,test_size=0.3,random_sta`

In [45]: `x_train.shape,x_test.shape,y_train.shape,y_test.shape`

Out[45]: ((1029, 33), (441, 33), (1029,), (441,))

Feature Scaling

```
In [46]: from sklearn.preprocessing import StandardScaler
```

```
In [47]: sc=StandardScaler()
```

```
In [48]: x_train=sc.fit_transform(x_train)
```

```
In [49]: x_test=sc.fit_transform(x_test)
```

Building the model

Multi-Linear Regression

```
In [50]: from sklearn.linear_model import LinearRegression
```

```
In [51]: lr = LinearRegression()
```

```
In [52]: lr.fit(x_train,y_train)
```

```
Out[52]: LinearRegression()
```

In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.

On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.

```
In [53]: lr.coef_ #slope(m)
```

```
Out[53]: array([-3.54940447e-02,  7.88352347e-05, -1.70825038e-02,  3.46389690e-02,
                2.44612841e-02,  3.65668214e-03,  4.16333634e-17, -9.46820520e-03,
                -4.11203734e-02,  1.06338881e-02, -2.97662154e-03, -3.84864283e-02,
                -1.52927977e-02, -1.57839139e-02, -3.67252862e-02,  3.35765928e-02,
                -5.90043558e-03,  5.81099165e-03,  3.78471890e-02, -6.93889390e-18,
                 9.55263279e-02, -2.55800078e-02,  2.01844797e-02, -2.64773510e-02,
                -1.21430643e-17, -1.79286106e-02, -3.30529386e-02, -1.09247807e-02,
                -3.10631611e-02, -2.47887717e-02, -1.10177742e-02,  2.11897289e-02,
                -6.60823991e-03])
```

```
In [54]: lr.intercept_ #(c)
```

```
Out[54]: 0.16229348882410102
```

```
In [55]: y_pred = lr.predict(x_test)
```


In [56]: `y_pred`

```
Out[56]: array([ 1.30302477e-01,  2.17626230e-01,  3.46282415e-01,  5.41382549e-03,
  4.99292896e-01,  1.01628868e-01,  3.44742777e-01,  1.23994945e-01,
 -1.60694945e-01,  4.02435622e-01,  1.44159172e-01,  2.67416840e-01,
 -4.62559536e-02,  5.58671849e-01,  2.81858700e-01,  1.53537792e-02,
  1.78573363e-01,  2.77532834e-01,  9.37121052e-02,  2.17571624e-01,
  2.65936178e-01,  1.41499184e-02,  8.36251186e-02,  9.58849826e-02,
  5.09869963e-01,  2.94764240e-01,  7.85819529e-02,  1.26647773e-01,
  5.05518902e-01,  8.48456917e-02, -7.97229275e-02,  2.15516993e-02,
  1.08079105e-01,  3.65998400e-01,  1.24517362e-01,  5.13682786e-02,
  1.06749689e-01,  6.07640778e-02,  6.66425313e-02,  4.81312859e-02,
 -1.16761425e-02, -2.97852924e-02,  5.25135582e-02, -1.59076817e-02,
 -1.71522795e-02,  4.17777714e-01,  3.67341564e-01, -2.14569245e-01,
  5.47964121e-01,  4.40723777e-01,  1.96701754e-01,  4.42415223e-01,
  1.45760263e-01,  3.75821843e-01,  4.92762622e-01,  2.95885645e-01,
 -4.62363391e-02,  3.16337190e-01, -7.90813313e-03,  2.52644685e-01,
 -3.18239329e-02,  2.83907645e-01,  9.03615010e-02,  1.26934391e-01,
  3.58670014e-01,  2.40923530e-02,  3.55890111e-01,  1.95961225e-01,
  1.28554515e-01,  1.18806226e-01, -2.86217094e-02,  3.17635336e-01,
  1.08017895e-01,  1.25723940e-01,  2.30183307e-01,  9.84315444e-02,
  0.10011060e-02,  0.70001435e-01,  0.50000700e-01,  0.00010750e-02])
```

In [57]: `y_test`

```
Out[57]: array([0, 0, 1, 0, 1, 0, 1, 0, 0, 1, 0, 1, 0, 1, 0, 0, 0, 1, 0, 0, 0, 0,
 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0,
 0, 1, 0, 0, 1, 1, 0, 1, 0, 0, 1, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0,
 1, 1, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0,
 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 1, 0, 0, 0, 0, 0, 0, 0,
 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0,
 1, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 1, 0, 0, 0, 0,
 0, 0, 0, 0, 0, 0, 1, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0,
 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 1, 0, 0, 1, 0, 0, 0, 0, 0,
 1, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 1, 0, 0, 0,
 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0,
 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 1, 0, 1, 0, 0, 0, 0,
 0, 0, 0, 1, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0,
 0, 1, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 1,
 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 1, 0, 0, 0, 0, 0, 0,
 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0,
 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 1, 0, 0, 1, 0, 0, 0, 0, 0,
 1, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 1, 0, 0, 1,
 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 1, 0, 0, 0, 0, 0,
 0, 0, 1, 0, 1, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0,
 0])
```

Logistic Regression

In [58]: `from sklearn.linear_model import LogisticRegression`

In [59]: `lg=LogisticRegression()`

```
In [60]: lg.fit(x_train,y_train)
```

```
Out[60]: LogisticRegression()
```

In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.

On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.

```
In [61]: y_pred_lg=lg.predict(x_test)
```

```
In [62]: y_pred
```

```
Out[62]: array([ 1.30302477e-01,  2.17626230e-01,  3.46282415e-01,  5.41382549e-03,
  4.99292896e-01,  1.01628868e-01,  3.44742777e-01,  1.23994945e-01,
 -1.60694945e-01,  4.02435622e-01,  1.44159172e-01,  2.67416840e-01,
 -4.62559536e-02,  5.58671849e-01,  2.81858700e-01,  1.53537792e-02,
  1.78573363e-01,  2.77532834e-01,  9.37121052e-02,  2.17571624e-01,
  2.65936178e-01,  1.41499184e-02,  8.36251186e-02,  9.58849826e-02,
  5.09869963e-01,  2.94764240e-01,  7.85819529e-02,  1.26647773e-01,
  5.05518902e-01,  8.48456917e-02, -7.97229275e-02,  2.15516993e-02,
  1.08079105e-01,  3.65998400e-01,  1.24517362e-01,  5.13682786e-02,
  1.06749689e-01,  6.07640778e-02,  6.66425313e-02,  4.81312859e-02,
 -1.16761425e-02, -2.97852924e-02,  5.25135582e-02, -1.59076817e-02,
 -1.71522795e-02,  4.17777714e-01,  3.67341564e-01, -2.14569245e-01,
  5.47964121e-01,  4.40723777e-01,  1.96701754e-01,  4.42415223e-01,
  1.45760263e-01,  3.75821843e-01,  4.92762622e-01,  2.95885645e-01,
 -4.62363391e-02,  3.16337190e-01, -7.90813313e-03,  2.52644685e-01,
 -3.18239329e-02,  2.83907645e-01,  9.03615010e-02,  1.26934391e-01,
  3.58670014e-01,  2.40923530e-02,  3.55890111e-01,  1.95961225e-01,
  1.28554515e-01,  1.18806226e-01, -2.86217094e-02,  3.17635336e-01,
  1.08017895e-01,  1.25723940e-01,  2.30183307e-01,  9.84315444e-02,
  0.10011000e-02,  0.70001425e-01,  0.50000700e-01,  4.00010750e-02])
```

```
Out[63]: array([[0, 0, 1, 0, 1, 0, 1, 0, 1, 0, 1, 0, 1, 0, 0, 1, 0, 0, 0, 0,
0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0,
0, 1, 0, 0, 1, 1, 0, 1, 0, 0, 1, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0,
1, 1, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0,
0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 1, 0, 0, 0, 0, 0, 0, 0,
0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 1,
1, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 1, 0, 0, 0, 0, 1,
0, 0, 0, 0, 0, 0, 1, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 1, 0,
0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 1, 0, 0, 1, 0, 0, 0, 0, 0,
1, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 1, 0, 0, 0, 0, 0,
0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0,
0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 1, 0, 1, 0, 0, 0, 0,
0, 0, 0, 1, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0,
0, 1, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 1,
0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 1, 0, 0, 0, 0, 0, 0,
0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0,
1, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 1, 0, 0, 1, 0,
0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 1, 0, 0, 0, 0, 0, 1,
0, 0, 1, 0, 1, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0,
0]])
```

```
In [64]: score = lg.score(x_test, y_test)
          print(score)
```

0.8820861678004536

Confusion matrix

```
In [65]: from sklearn import metrics
cm = metrics.confusion_matrix(y_test,y_pred_lg)
print(cm)
```

$$\begin{bmatrix} 366 & 5 \\ 47 & 23 \end{bmatrix}$$

Ridge and Lasso

```
In [66]: from sklearn.linear_model import Ridge
         from sklearn.model_selection import GridSearchCV
```

```
In [67]: rg=Ridge()
```

```
In [68]: parameters={"alpha": [1, 2, 3, 5, 10, 20, 30, 40, 60, 70, 80, 90]}
ridgecv=GridSearchCV(rg, parameters, scoring="neg_mean_squared_error", cv=5)
ridgecv.fit(x_train, y_train)
```

```
Out[68]: GridSearchCV(cv=5, estimator=Ridge(),
                    param_grid={'alpha': [1, 2, 3, 5, 10, 20, 30, 40, 60, 70, 80, 90]},
                    scoring='neg_mean_squared_error')
```

In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.

On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.

```
In [69]: print(ridgecv.best_params_)

{'alpha': 90}
```

```
In [70]: print(ridgecv.best_score_)

-0.11390621139234183
```

```
In [71]: y_pred_rg=ridgecv.predict(x_test)
```

```
In [72]: y_pred_rg
```

```
Out[72]: array([ 1.34413485e-01,  2.22561818e-01,  3.41692977e-01,  3.88209867e-03,
  4.84617338e-01,  1.16361483e-01,  3.30449743e-01,  1.27358807e-01,
 -1.34442619e-01,  3.77692888e-01,  1.33001445e-01,  2.69898751e-01,
 -2.54707392e-02,  5.25771894e-01,  2.67543514e-01,  2.78725024e-02,
  1.82233111e-01,  2.78896415e-01,  9.12689699e-02,  2.11494641e-01,
  2.70103341e-01,  8.44922044e-03,  8.74746722e-02,  1.05348798e-01,
  4.87749940e-01,  2.83080512e-01,  8.80556209e-02,  1.23817268e-01,
  4.82185624e-01,  9.34824523e-02, -7.16448509e-02,  4.07003104e-02,
  1.08437994e-01,  3.42151399e-01,  1.22270929e-01,  6.85889862e-02,
  1.06690533e-01,  7.08689637e-02,  7.51570276e-02,  6.05829413e-02,
  1.08782897e-02, -6.91368661e-03,  5.83191600e-02, -1.54680056e-02,
 -4.02267475e-03,  4.08010612e-01,  3.43668700e-01, -1.83519405e-01,
  5.29536511e-01,  4.27646098e-01,  1.95234877e-01,  4.25012930e-01,
  1.40754410e-01,  3.52173952e-01,  4.70372694e-01,  2.89240343e-01,
 -3.11642726e-02,  3.04206456e-01,  9.89337674e-03,  2.44569884e-01,
 -1.40249115e-02,  2.75133912e-01,  8.64669565e-02,  1.24214885e-01,
  3.48994545e-01,  3.41026778e-02,  3.40548051e-01,  1.95847356e-01,
  1.30040885e-01,  1.32259137e-01, -2.34680143e-02,  3.04595468e-01,
  1.12452197e-01,  1.30525275e-01,  2.19329505e-01,  9.44722098e-02,
  0.00000000e+00,  0.00000000e+00,  0.00000000e+00,  0.00000000e+00])
```

```
Out[73]: array([[0, 0, 1, 0, 1, 0, 1, 0, 0, 1, 0, 1, 0, 1, 0, 0, 0, 1, 0, 0, 0, 0,
0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0,
0, 1, 0, 0, 1, 1, 0, 1, 0, 0, 1, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0,
1, 1, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0,
0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 1, 0, 0, 0, 0, 0, 0, 0, 0,
0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 1,
1, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 1, 0, 0, 0, 0, 1,
0, 0, 0, 0, 0, 0, 1, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 1, 0,
0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 1, 0, 0, 1, 0, 0, 0, 0, 0, 0,
1, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 1, 0, 0, 0, 0, 0,
0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0,
0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 1, 0, 1, 0, 0, 0, 0,
0, 0, 0, 1, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0,
0, 1, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 1,
0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 1, 0, 0, 0, 0, 0, 0,
0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 1, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0,
1, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 1, 0, 0, 1, 0,
0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 1, 0, 0, 0, 0, 0, 1,
0, 0, 1, 0, 1, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0,
0]])
```

Lasso

```
In [77]: parameters={"alpha":[1,2,3,5,10,20,30,40,60,70,80,90]}
         ridgecv=GridSearchCV(la,parameters,scoring="neg_mean_squared_error",cv=5)
         ridgecv.fit(x_train,y_train)
```

In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.

On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.

```
In [78]: print(ridgecv.best_params_)
```

```
{'alpha': 90}
```

```
In [79]: print(ridgecv.best_score_)
```

```
-0.11390621139234183
```

```
In [80]: y_pred_la=ridgecv.predict(x_test)
```

```
In [81]: y_pred_la
```

```
Out[81]: array([ 1.34413485e-01,  2.22561818e-01,  3.41692977e-01,  3.88209867e-03,
  4.84617338e-01,  1.16361483e-01,  3.30449743e-01,  1.27358807e-01,
 -1.34442619e-01,  3.77692888e-01,  1.33001445e-01,  2.69898751e-01,
 -2.54707392e-02,  5.25771894e-01,  2.67543514e-01,  2.78725024e-02,
  1.82233111e-01,  2.78896415e-01,  9.12689699e-02,  2.11494641e-01,
  2.70103341e-01,  8.44922044e-03,  8.74746722e-02,  1.05348798e-01,
  4.87749940e-01,  2.83080512e-01,  8.80556209e-02,  1.23817268e-01,
  4.82185624e-01,  9.34824523e-02, -7.16448509e-02,  4.07003104e-02,
  1.08437994e-01,  3.42151399e-01,  1.22270929e-01,  6.85889862e-02,
  1.06690533e-01,  7.08689637e-02,  7.51570276e-02,  6.05829413e-02,
  1.08782897e-02, -6.91368661e-03,  5.83191600e-02, -1.54680056e-02,
 -4.02267475e-03,  4.08010612e-01,  3.43668700e-01, -1.83519405e-01,
  5.29536511e-01,  4.27646098e-01,  1.95234877e-01,  4.25012930e-01,
  1.40754410e-01,  3.52173952e-01,  4.70372694e-01,  2.89240343e-01,
 -3.11642726e-02,  3.04206456e-01,  9.89337674e-03,  2.44569884e-01,
 -1.40249115e-02,  2.75133912e-01,  8.64669565e-02,  1.24214885e-01,
  3.48994545e-01,  3.41026778e-02,  3.40548051e-01,  1.95847356e-01,
  1.30040885e-01,  1.32259137e-01, -2.34680143e-02,  3.04595468e-01,
  1.12452197e-01,  1.30525275e-01,  2.19329505e-01,  9.44722098e-02,
  0.00105700e-02,  0.00105700e-02,  0.00105700e-02,  0.00105700e-02])
```

```
In [82]: from sklearn import metrics
print(metrics.r2_score(y_test,y_pred_la))
print(metrics.r2_score(y_train,ridgecv.predict(x_train)))
```

```
0.21073458438815906
```

```
0.2061567210285109
```

Decision Tree

```
In [83]: from sklearn.tree import DecisionTreeClassifier
dtc=DecisionTreeClassifier()
```

```
In [84]: dtc.fit(x_train,y_train)
```

```
Out[84]: DecisionTreeClassifier()
```

In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.

On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.

```
In [85]: pred=dtc.predict(x_test)
```

```
In [86]: pred
```

```
Out[86]: array([0, 0, 0, 0, 1, 1, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 1, 0, 0, 0, 0,
                0, 0, 0, 0, 0, 0, 1, 0, 1, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
                0, 0, 0, 0, 1, 0, 0, 0, 1, 0, 1, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0,
                0, 0, 0, 0, 1, 0, 0, 0, 1, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
                0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0,
                0, 1, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1,
                1, 1, 0, 1, 0, 1, 0, 1, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0,
                0, 0, 1, 0, 0, 0, 1, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
                0, 1, 0, 1, 0, 0, 1, 1, 0, 0, 0, 0, 0, 1, 0, 1, 0, 0, 0, 0, 0, 0,
                0, 0, 0, 0, 1, 0, 1, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1,
                0, 1, 1, 0, 0, 0, 0, 1, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
                0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 1, 0, 0, 1, 0,
                0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 1, 0, 0, 0, 0, 1,
                0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 1, 0, 1, 0, 0, 0, 0, 0, 0,
                0, 1, 1, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 1, 1, 0, 1, 1, 0, 0,
                0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 1, 0, 1, 0, 0, 0, 0, 0, 1, 0, 1,
                0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
                0, 0, 0, 0, 1, 0, 1, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 1, 0, 1, 0,
                0, 1, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 1, 1, 0, 0, 0, 0,
                0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0,
                0])
```

```
In [87]: y_test
```

```
Out[87]: array([0, 0, 1, 0, 1, 0, 1, 0, 0, 1, 0, 1, 0, 1, 0, 0, 0, 1, 0, 0, 0, 0,
                0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0,
                0, 1, 0, 0, 1, 1, 0, 1, 0, 0, 1, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0,
                1, 1, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0,
                0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 1, 0, 0, 0, 0, 0, 0, 0,
                0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 1,
                1, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 1, 0, 0, 0, 1,
                0, 0, 0, 0, 0, 0, 1, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 1, 0,
                0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 1, 0, 0, 1, 0, 0, 0, 0, 0,
                1, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 1, 0, 0, 0, 0,
                0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0,
                0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 1, 0, 1, 0, 0, 0, 0,
                0, 0, 0, 1, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0,
                0, 1, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 1,
                0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 1, 0, 0, 0, 0, 0,
                0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0,
                0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 1, 0, 0, 1, 0, 0, 0, 0, 0,
                1, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 1, 0, 0, 1, 0,
                0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 1, 0, 0, 0, 0, 0, 1,
                0, 0, 1, 0, 1, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0,
                0])
```

```
In [88]: #Accuracy score
from sklearn.metrics import accuracy_score, confusion_matrix, classification_rep
```

```
In [89]: accuracy_score(y_test,pred)
```

```
Out[89]: 0.7664399092970522
```

```
In [90]: confusion_matrix(y_test,pred)
```

```
Out[90]: array([[316,  55],
                [ 48,  22]], dtype=int64)
```

```
In [91]: pd.crosstab(y_test,pred)
```

```
Out[91]:
```

	col_0	0	1
row_0			
0	316	55	
1	48	22	

```
In [92]: print(classification_report(y_test,pred))
```


	precision	recall	f1-score	support
0	0.87	0.85	0.86	371
1	0.29	0.31	0.30	70
accuracy			0.77	441
macro avg	0.58	0.58	0.58	441
weighted avg	0.78	0.77	0.77	441

```
In [93]: probability=dtc.predict_proba(x_test)[:,-1]
```

```
In [94]: # roc_curve
fpr,tpr,threshholds = roc_curve(y_test,probability)
```



```
In [95]: plt.plot(fpr, tpr)
plt.xlabel('FPR')
plt.ylabel('TPR')
plt.title('ROC CURVE')
plt.show()
```



Random Forest

```
In [96]: from sklearn.ensemble import RandomForestClassifier
rfc=RandomForestClassifier()
```

```
In [97]: forest_params = [{'max_depth': list(range(10, 15)), 'max_features': list(range
```

```
In [98]: from sklearn.model_selection import GridSearchCV
```

```
In [99]: rfc_cv= GridSearchCV(rfc,param_grid=forest_params,cv=10,scoring="accuracy")
```

```
In [100]: rfc_cv.fit(x_train,y_train)
```

```
C:\Users\Prasanth Nimmala\anaconda3\lib\site-packages\sklearn\model_selection\_validation.py:378: FitFailedWarning:
50 fits failed out of a total of 700.
The score on these train-test partitions for these parameters will be set to nan.
If these failures are not expected, you can try to debug them by setting error_score='raise'.
```

```
Below are more details about the failures:
```

```
-----
---
50 fits failed with the following error:
Traceback (most recent call last):
  File "C:\Users\Prasanth Nimmala\anaconda3\lib\site-packages\sklearn\model_selection\_validation.py", line 686, in _fit_and_score
    estimator.fit(X_train, y_train, **fit_params)
  File "C:\Users\Prasanth Nimmala\anaconda3\lib\site-packages\sklearn\ensemble\_forest.py", line 340, in fit
    self._validate_params()
  File "C:\Users\Prasanth Nimmala\anaconda3\lib\site-packages\sklearn\base.py", line 581, in _validate_params
    validate_parameter_constraints(
  File "C:\Users\Prasanth Nimmala\anaconda3\lib\site-packages\sklearn\utils\_param_validation.py", line 97, in validate_parameter_constraints
    raise InvalidParameterError(
sklearn.utils._param_validation.InvalidParameterError: The 'max_features' parameter of RandomForestClassifier must be an int in the range [1, inf), a float in the range (0.0, 1.0], a str among {'log2', 'sqrt', 'auto' (deprecated)} or None. Got 0 instead.
```

```
warnings.warn(some_fits_failed_message, FitFailedWarning)
C:\Users\Prasanth Nimmala\anaconda3\lib\site-packages\sklearn\model_selection\_search.py:952: UserWarning: One or more of the test scores are non-finite:
[      nan  0.84353703  0.84840091  0.8483914   0.85325528  0.85033314
  0.85421664  0.85033314  0.85422616  0.84644013  0.85517799  0.85519703
  0.85033314  0.84449838      nan  0.8445079   0.84935275  0.85031411
  0.85421664  0.84936227  0.85516848  0.85032362  0.84934323  0.8512945
  0.84935275  0.84934323  0.85322673  0.85032362      nan  0.8445079
  0.84936227  0.85324576  0.85033314  0.85033314  0.85324576  0.85810013
  0.85711974  0.84935275  0.85225585  0.8483914   0.85131354  0.85324576
      nan  0.84546926  0.84937179  0.84936227  0.85325528  0.85324576
  0.85615839  0.85324576  0.85520655  0.85615839  0.85517799  0.85324576
  0.8512945   0.85030459      nan  0.84547877  0.84644965  0.84546926
  0.85518751  0.84353703  0.84937179  0.85615839  0.85031411  0.8561679
  0.85713878  0.84838188  0.85227489  0.84643061]
warnings.warn(
```

```
Out[100]: GridSearchCV(cv=10, estimator=RandomForestClassifier(),
                      param_grid=[{'max_depth': [10, 11, 12, 13, 14],
                                   'max_features': [0, 1, 2, 3, 4, 5, 6, 7, 8, 9, 10,
11,
                                                12, 13]}],
                      scoring='accuracy')
```

In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.

On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.

```
In [101]: pred=rfc_cv.predict(x_test)
```

```
In [102]: print(classification_report(y_test,pred))
```

	precision	recall	f1-score	support
0	0.87	0.99	0.93	371
1	0.84	0.23	0.36	70
accuracy			0.87	441
macro avg	0.86	0.61	0.64	441
weighted avg	0.87	0.87	0.84	441

```
In [103]: rfc_cv.best_params_
```

```
Out[103]: {'max_depth': 12, 'max_features': 7}
```

```
In [104]: rfc_cv.best_score_
```

```
Out[104]: 0.8581001332571864
```